

## EL - bagging (pima-indians-diabetes)

December 18, 2022

```
[3]: import pandas as pd
```

```
df = pd.read_csv("pima-indians-diabetes.csv")  
df.head()
```

```
[3]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

  

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[4]: df.isnull().sum()
```

```
[4]: Pregnancies      0  
Glucose      0  
BloodPressure  0  
SkinThickness  0  
Insulin      0  
BMI          0  
DiabetesPedigreeFunction  0  
Age          0  
Outcome      0  
dtype: int64
```

```
[5]: df.describe()
```

```
[5]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	\
count	768.000000	768.000000	768.000000	768.000000	768.000000	
mean	3.845052	120.894531	69.105469	20.536458	79.799479	

std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
[6]: df.Outcome.value_counts()
```

```
[6]: 0    500
      1    268
      Name: Outcome, dtype: int64
```

```
[7]: X = df.drop("Outcome",axis="columns")
      y = df.Outcome
```

```
[8]: from sklearn.preprocessing import StandardScaler

      scaler = StandardScaler()
      X_scaled = scaler.fit_transform(X)
      X_scaled[:3]
```

```
[8]: array([[ 0.63994726,  0.84832379,  0.14964075,  0.90726993, -0.69289057,
              0.20401277,  0.46849198,  1.4259954 ],
            [-0.84488505, -1.12339636, -0.16054575,  0.53090156, -0.69289057,
              -0.68442195, -0.36506078, -0.19067191],
            [ 1.23388019,  1.94372388, -0.26394125, -1.28821221, -0.69289057,
              -1.10325546,  0.60439732, -0.10558415]])
```

```
[9]: from sklearn.model_selection import train_test_split

      X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, stratify=y,
      ↪random_state=10)
```

```
[10]: X_train.shape
```

```
[10]: (576, 8)
```

```
[13]: X_test.shape
```

```
[13]: (192, 8)
```

```
[14]: y_train.value_counts()
```

```
[14]: 0    375  
      1    201  
      Name: Outcome, dtype: int64
```

```
[15]: 201/375
```

```
[15]: 0.536
```

```
[16]: y_test.value_counts()
```

```
[16]: 0    125  
      1     67  
      Name: Outcome, dtype: int64
```

```
[17]: 67/125
```

```
[17]: 0.536
```

```
[18]: from sklearn.model_selection import cross_val_score  
      from sklearn.tree import DecisionTreeClassifier  
  
      scores = cross_val_score(DecisionTreeClassifier(), X, y, cv=5)  
      scores
```

```
[18]: array([0.68831169, 0.65584416, 0.67532468, 0.78431373, 0.71895425])
```

```
[19]: scores.mean()
```

```
[19]: 0.7045496986673457
```

```
[20]: from sklearn.ensemble import BaggingClassifier  
  
      bag_model = BaggingClassifier(  
          base_estimator=DecisionTreeClassifier(),  
          n_estimators=100,  
          max_samples=0.8,  
          oob_score=True,  
          random_state=0  
      )  
      bag_model.fit(X_train, y_train)  
      bag_model.oob_score_
```

[20]: 0.7534722222222222

```
[21]: bag_model.score(X_test, y_test)
```

[21]: 0.7760416666666666

```
[22]: bag_model = BaggingClassifier(  
    base_estimator=DecisionTreeClassifier(),  
    n_estimators=100,  
    max_samples=0.8,  
    oob_score=True,  
    random_state=0  
)  
scores = cross_val_score(bag_model, X, y, cv=5)  
scores.mean()
```

[22]: 0.7578728461081402

```
[40]: from sklearn.ensemble import RandomForestClassifier  
  
scores = cross_val_score(RandomForestClassifier(n_estimators=50), X, y, cv=5)  
scores.mean()
```

[40]: 0.7747899159663865

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```